

Artificial Analysis Coding Agent Benchmarks

We measure real-world performance of coding agents on software engineering tasks, including cost, token usage, and execution time. We compare how performance changes across agents, models, and execution settings.

To compare language models see our [model benchmarks](#).

Benchmarks  Comparisons  Features

Artificial Analysis Coding Agent Index


Composite index of 3 benchmarks:


- DeepSWE
Software engineering tasks, 113 tasks
[By Datacurve](#)
- Terminal-Bench v2
Agentic terminal use, 84 tasks
[By Laude Institute](#)
- SWE-Atlas-QnA
Technical Q&A, 124 tasks
[By Scale AI](#)

Index represents the average pass@1 across 3 runs of each benchmark. Index recently updated to v1.1. [See methodology for details](#)

 [Coding Agents](#)  General Work  Chatbots  Presentations  OCR


Highlights






Time per Task

Average agent wall time per task · Lower is better




Cost per Task

Average API cost per task (USD) · Lower is better



Performance

Performance across the Artificial Analysis Coding Agent Index.




Harness Comparison: Artificial Analysis Coding Agent Index

Composite average pass@1 across Claude Code, Cursor CLI, and Opencode for Claude Opus 4.7 · Higher is better



Artificial Analysis



What This Chart Shows

+

Token Usage

Token consumption across the Artificial Analysis Coding Agent Index, including total usage, token mix, efficiency, and per-benchmark breakdowns.

Token Usage per Task


Average input, cache, and output tokens per task



14 of 43 models

Output Cached Input Input

Artificial Analysis



ⓘ Prompt cache hit rates can vary significantly by provider routing, which can materially change effective cost.

- ⓘ Input Tokens +
- ⓘ Cached Input Tokens +
- ⓘ Why Cache Usage Varies +
- ⓘ Output Tokens +

Artificial Analysis Coding Agent Index vs. Total Tokens


Artificial Analysis Coding Agent Index vs. average total tokens per task



Color by Model Agent 14 of 43 models

Most attractive quadrant

OpenAI Anthropic SpaceXAI Meta Z.ai Cursor DeepSeek Google



How to Read This Chart

Artificial Analysis

Premium

Log in

Q

≡

Cost

Cost across the Artificial Analysis Coding Agent Index based on current per-token API pricing, including cache write pricing and cache discounts where available. Many users will access coding agent harnesses through subscription plan offerings rather than pay-per-token.

Cost to Run

Cost Distribution

Total Cost

Cost per Task

Average pay-per-token API cost per task (USD) · Lower is better


Color by

Model

Agent

14 of 43 models





What Cost Is Measuring +

Artificial Analysis Coding Agent Index vs. Cost per Task


Artificial Analysis Coding Agent Index vs. average pay-per-token API cost per task (USD)



Color by **Model** **Agent** 14 of 43 models

☐ Most attractive quadrant

OpenAI Anthropic SpaceXAI Meta Z.ai Cursor DeepSeek Google



How to Read This Chart



Execution Time

Active agent runtime across the Artificial Analysis Coding Agent Index.

Execution Time

Turns

Time per Task

Average agent wall time per task · Lower is better




Color by

Model

Agent

14 of 43 models





i What Execution Time Is Measuring +

Artificial Analysis Coding Agent Index vs. Execution Time


Artificial Analysis Coding Agent Index vs. average agent wall time per task



Color by **Model** **Agent** 14 of 43 models ^

☐ Most attractive quadrant

OpenAI Anthropic SpaceXAI Meta Z.ai Cursor DeepSeek Google



How to Read This Chart

+

Run Specifications

+

Frequently Asked Questions

What is the Artificial Analysis Coding Agent Index?

+

Which benchmarks are included in the index right now?

+

What kinds of tasks are these benchmarks actually testing?

+

How do Q&A-style tasks differ from implementation-style tasks?

+

How are agents scored on each benchmark?



How is the overall index weighted?



What does execution time mean?




What does token usage mean, and why does it matter?



Why can a higher-index agent still be worse for my use case?



How realistic are these tasks, and what setup was used for each agent?



Artificial Analysis

Get notified about new articles

What's your email?

Subscribe

Explore

Company

LLM Leaderboard

Image Arena

Video Arena

AI Agents

Evaluations

Services

Contact

Articles

FAQ

